

# Real-Time Marginal Ambiguity Covariance for Tightly Coupled GNSS–Visual–Inertial Factor Graphs

Anonymous Authors

**Abstract**—Carrier-phase ambiguity resolution (AR) in tightly coupled GNSS/IMU/vision estimators consumes a float ambiguity covariance  $Q_{aa}$ . Open platforms often substitute a GNSS-only “shadow” covariance because forming the joint sliding-window covariance online is considered too expensive, discarding the IMU and visual information that tight coupling is meant to provide. We present an exact-in-window marginal at the optimizer’s fixed linearization,  $Q_{aa} = [H_{\text{active}}^{-1}]_{aa}$  obtained by two-stage sparse Schur elimination over the active GNSS–IMU–vision window, including the marginalization prior, with a numerical self-validation gate and shadow fallback. On 1 787 open-sky AR epochs the method produces a validated covariance on 99.94% of epochs, with 90.1% achieving relative Frobenius error  $\|Q_{\text{prop}} - Q_{\text{Ceres}}\|_F / \|Q_{\text{Ceres}}\|_F$  below  $10^{-3}$ ; among the remaining disagreements a trace proxy is larger (conservative) in 171 of 177 cases and within  $7.3 \times 10^{-4}$  in 6 near-ties. Mean covariance query time is 28 ms versus 812 ms for Ceres ( $\sim 29\times$ ), with median 27 ms, p95 49 ms, p99 59 ms, and 4.3% of epochs above a 50 ms covariance budget (max 97.5 ms). End-to-end effects versus shadow AR are reported as secondary system impact; an overconfident local-conditioning approximation is harmful. Directional Loewner diagnostics and a Ceres-fed end-to-end AR oracle are left to future validation; accordingly, we claim neither directional dominance nor trajectory equivalence to Ceres (Sec. IX).

**Index Terms**—RTK, ambiguity resolution, factor graph, marginal covariance, GNSS/IMU/vision, Schur complement

## I. Introduction

Tightly coupled GNSS/IMU/vision estimators resolve carrier-phase integers using a float ambiguity covariance fed to LAMBDA-type search and acceptance tests. In GICI-LIB [1], the open RTK/IMU/camera reference, AR consumes a covariance from a parallel GNSS-only shadow estimator because a full-graph query via `ceres::Covariance` is too slow online (mean 812 ms on our filled-window open-sky graphs; 99.6% of epochs exceed 50 ms). The shadow path discards IMU and visual information at the moment AR needs it most.

Cheaper alternatives that condition on neighboring states without forming the Schur complement are fast but can be overconfident: on UrbanNav Deep, a local-conditioning variant raised fix rate while degrading horizontal RMSE from  $2.77 \pm 0.44$  m to  $6.75 \pm 2.98$  m. AR therefore needs a covariance that is both consistent with the joint window and affordable every epoch.

Manuscript prepared for IEEE Robotics and Automation Letters (RAL). Numerical claims are traced to logged per-epoch benchmarks and results/research/paper\_repro/ where cited. Ground truth is used only in post-hoc evaluation.

This paper focuses on that covariance problem. We make three contributions:

- 1) An algorithm that extracts the exact-in-window, fixed-linearization marginal ambiguity covariance  $Q_{aa} = [H_{\text{active}}^{-1}]_{aa}$  by structure-exploiting two-stage elimination, without inverting the full information matrix and without dropping the marginalization prior.
- 2) A numerical self-validation gate with typed abstention and shadow fallback, so an unreliable factorization never feeds AR.
- 3) A covariance-level evaluation against shadow AR, local conditioning, and `ceres::Covariance`, with end-to-end system impact reported secondarily.

The two evaluation questions are: (i) does the proposed path reproduce the Ceres ambiguity covariance (and, offline, Ceres-driven AR decisions)? (ii) by how much does it reduce cost relative to the estimator/AR cycle? Trajectory RMSE is secondary and is not the headline claim.

## II. Related Work

a) Integer estimation.: ILS via LAMBDA [2] / MLAMBDA [3], ratio tests and FFRT [4], [5], partial AR, and bootstrapping success rates  $P_s$  [6] are standard. BIE [7], [8] softens hard commits. Our focus is the float covariance that those methods consume inside a multi-sensor graph, not a new integer estimator.

b) GNSS/IMU/vision fusion.: Tight VIO [9], [10] and GNSS–VIO systems [1], [11] optimize sliding-window factor graphs with marginalization priors. GICI-LIB’s AR path uses a GNSS-only shadow covariance—the gap we close at the covariance interface.

c) Covariance recovery in factor graphs.: Marginal covariances are classical in Gauss–Newton estimation. Sparse factorization and selected inversion recover subsets of  $H^{-1}$  without forming the full dense inverse [12], [13]. Incremental smoothers (iSAM2 / Bayes tree) maintain square-root information and can extract covariances from the clique structure [14], [15]. Ceres exposes sparse-QR covariance queries [16] that are exact for the linearized problem but costly on large live windows. Our method is a domain-structured elimination for the ambiguity block of a GNSS/IMU/vision RRR window: landmark Schur first, then dense reduction of the remaining nuisance states, with an AR-oriented abstention policy.




Fig. 1. RRR window and AR path. The shaded block is the proposed exact-in-window marginal; the baseline uses a GNSS-only shadow covariance instead.

d) Urban GNSS.: NLOS/multipath mitigations include masks, robust losses, shadow matching [17], and 3DMA [18], [19]. UrbanNav-HK [20] provides Medium/Deep benchmarks. We use these datasets to stress end-to-end AR under the proposed  $Q_{aa}$ ; we do not claim a new NLOS classifier.

### III. Problem Formulation

We consider a Ceres-based sliding-window RRR graph with double-differenced GNSS residuals, IMU preintegration, reprojection factors, and a dense marginalization prior. Let  $x$  collect active states and  $H = \sum_r J_r^\top J_r$  the Gauss-Newton information in minimal coordinates (Triggs-corrected Jacobians for robust residuals), including the prior information  $\Lambda$ . Partition  $x = (a, \ell, n)$  into ambiguities  $a$ , landmarks  $\ell$ , and remaining nuisance  $n$  (poses, speed/bias, clocks, frequencies, extrinsics, prior-only variables).

AR requires the marginal

$$Q_{aa} = [H^{-1}]_{aa}. \quad (1)$$

ceres::Covariance computes (1) at whole-problem sparse-QR cost. The shadow estimator replaces (1) with a GNSS-only approximation. A local-conditioning shortcut drops the Schur term involving out-of-set neighbors and is fast but can understate uncertainty.

a) Requirements.: (R1) Exactness in the active window:  $Q_{aa}$  must equal the Ceres marginal up to linearization and numerics. (R2) Real-time: cost bounded by window size, with a fallback when factorization is unreliable. (R3) AR safety preference: when numerics disagree, prefer larger uncertainty over silent overconfidence.

Fig. 1 places  $Q_{aa}$  in the AR path.

### IV. Proposed Marginal Covariance Method

a) Assembly.: We form  $H$  from all active residuals plus the marginalization prior through a read-only accessor exposing  $\Lambda = J^\top J$  with per-block minimal offsets (generic residual evaluation of the prior is avoided under live window growth).

---

#### Algorithm 1 Exact-in-window marginal ambiguity covariance

---

Require: active residuals, prior  $\Lambda$ , ambiguity index set  $a$

- 1: Assemble  $H \leftarrow \sum_r J_r^\top J_r$  (Triggs)  $+\Lambda$
- 2: for each landmark block  $\ell$  do
- 3:   Rank-truncated Schur-eliminate  $\ell$  from  $H$
- 4: end for
- 5: Jacobi-equilibrate nuisance; LDLT factorize
- 6: Iterative refine; if gate fails then return shadow  $Q$
- 7: return  $Q_{aa} \leftarrow [H_{\text{red}}^{-1}]_{aa}$

---

b) Two-stage elimination.: No residual couples an ambiguity to a landmark ( $H_{a\ell} = 0$ ; prior-touched landmarks are promoted into  $n$ ), so

$$H = \begin{bmatrix} H_{aa} & H_{an} & 0 \\ H_{na} & H_{nn} & H_{n\ell} \\ 0 & H_{\ell n} & H_{\ell\ell} \end{bmatrix}. \quad (2)$$

Stage 1 forms the landmark Schur update

$$\bar{H}_{nn} = H_{nn} - H_{n\ell} H_{\ell\ell}^+ H_{\ell n} \quad (3)$$

via independent  $3 \times 3$  rank-truncated pseudo-inverses. Exactness condition: null modes of each  $H_{ii}$  have zero coupling rows in  $H_{*i}$ , so truncating those modes does not change the  $(a, n)$  marginal. Stage 2 then yields the ambiguity marginal

$$Q_{aa} = \left( H_{aa} - H_{an} \bar{H}_{nn}^{-1} H_{na} \right)^{-1} \quad (4)$$

by Jacobi-equilibrated dense LDLT on  $\bar{H}_{nn}$  plus iterative refinement (Fig. 2, Alg. 1).

c) Relation to standard covariance recovery.: Ceres selected covariance [16] asks for the same blocks but factors the whole problem (sparse QR) each query. Sparse selected inversion / Takahashi methods recover entries of  $H^{-1}$  from a global factorization [12], [13]; Bayes-tree methods reuse incremental square-root information [14]. Our contribution is not a new linear-algebra primitive but a domain-structured online path: landmark blocks are eliminated at  $O(1)$  each, the remaining GNSS/IMU/extrinsic/prior nuisance is small enough for dense reduction every AR epoch, the marginalization prior is included through a read-only  $\Lambda$  accessor, and a typed abstention prevents feeding AR an untrusted  $Q_{aa}$ . Complexity:  $O(L)$  for  $L$  landmarks ( $3 \times 3$  each) plus  $O(n_n^3 + n_a n_n^2)$  for nuisance size  $n_n$  and ambiguity size  $n_a$ , bounded by the sliding window.

d) Self-validation and fallback.: After LDLT, a few refinement steps on the reduced system are accepted only if the last increment's effect on the Schur complement is below the scale of the smallest eigenvalue of the reduced nuisance block (no dataset-tuned thresholds). Abstention cases: non-finite assembly, failed factorization, failed gate, or non-PD reduced ambiguity information. On abstention the estimator returns the shadow covariance for that epoch.




Fig. 2. Two-stage elimination with abstention (F2).

TABLE I  
Covariance vs. ceres::Covariance (GICI-board 1.1)

| Metric                              | Value                           |
|-------------------------------------|---------------------------------|
| Validated output (fast_ok)          | 99.94% of 1787 (1 abstention)   |
| Frobenius rel. err. $\leq 10^{-3}$  | 90.1% of usable                 |
| Frobenius disagreements $> 10^{-3}$ | 177                             |
| trace larger (proxy)                | 171                             |
| trace below-ref. (near-tie)         | 6 (worst $7.3 \times 10^{-4}$ ) |
| ASan (1761 epochs)                  | clean                           |

The estimator is registered as an opt-in variant (rtk\_imu\_camera\_rrr\_va); upstream defaults are unchanged. File-mode UrbanNav configs disclose the retained max\_age=35s setting and relative\_frequency=1.0; neither was tuned against UrbanNav ground truth.

## V. Numerical Validation and Complexity


a) Reference.: Per-AR-epoch comparison against ceres::Covariance on identical live graphs (GICI-board 1.1, 1787 epochs), requesting the same ambiguity blocks (not a whole-state inverse straw man). The primary match metric already logged online is the normalized Frobenius error  $\|Q_{\text{prop}} - Q_{\text{Ceres}}\|_F / \|Q_{\text{Ceres}}\|_F$ . Trace comparisons are a secondary safety proxy only. Loewner / directional diagnostics (min eig( $Q_{\text{prop}} - Q_{\text{Ceres}}$ ), generalized eigenvalues of the pair) are outside the evidence reported here; we therefore make no directional-dominance claim (Sec. IX).

b) Accuracy.: Table I: a validated covariance is produced on 99.94% of epochs (one abstention)—this is not a 99.94% match rate. Of usable epochs, 90.1% have Frobenius relative error  $\leq 10^{-3}$ . Of 177 larger Frobenius disagreements, a trace proxy is larger in 171 cases and within  $7.3 \times 10^{-4}$  in 6 near-ties. Larger trace does not imply Loewner dominance in every direction.

c) Runtime.: Table II and Fig. 3 report the covariance query only. AR in this stack is invoked near GNSS epoch rate ( $\sim 1$  Hz); a 50 ms covariance budget is a soft online target ( $\sim 5\%$  of a 1 Hz slot), not a hard 20 Hz estimator deadline. Mean proposed query 28 ms ( $\sim 29\times$  vs Ceres); 4.3% exceed 50 ms (max 97.5 ms). Total optimize+covariance latency will be reported with confirmatory runs.

TABLE II  
Runtime on the same 1787 AR epochs (ms)

| Method   | mean  | med   | p95  | p99  | max  | > 50 ms |
|----------|-------|-------|------|------|------|---------|
| Proposed | 28.0  | 26.7  | 49.1 | 58.8 | 97.5 | 4.3%    |
| Ceres    | 811.9 | 758.1 | 1573 | 1892 | 2714 | 99.6%   |

Fig. 3. Proposed vs. Ceres wall time versus graph size (GICI-board 1.1,  $n = 1787$ ).TABLE III  
Datasets

| Dataset        | Role           | DOY         | Notes     |
|----------------|----------------|-------------|-----------|
| GICI-board 1.1 | Cov. + runtime | —           | Open-sky  |
| UrbanNav Deep  | System impact  | 141 Whampoa | $n=15119$ |

d) Complexity.: Stage 1 is linear in the number of landmarks with  $3 \times 3$  work each. Stage 2 is cubic in the dense nuisance dimension after landmark elimination, hence bounded by the sliding window rather than trajectory length—the same structural reason the method tracks Ceres' marginal while remaining cheaper than a whole-problem sparse QR each epoch.

## VI. Experimental Setup

UrbanNav Deep uses the session-aligned 5 s HKKT base slice selected by the repository runner (not the full-day 30 s file). Primary metric: unaligned raw ENU RMSE (E/N/U, yaw). Sim(3) APE is used only when comparing to Chi Table V. Admissibility: Deep  $n_{\text{matched}} \geq 14000$ . Research flags are opt-in and default-off; the frozen upstream baseline APE on GICI-board 1.1 (0.029198 m / 0.470557°, bands  $\pm 0.005$  m /  $\pm 0.1^\circ$ ) is re-checked after builds.

Post-file replay is load-sensitive (same Deep sequence spans  $\sim 14$ – $24$  min wall under identical configs). We treat this as a reproducibility limitation (Sec. IX), not a scientific contribution; confirmatory UrbanNav numbers below use identical-load pairing or solo-idle execution.

TABLE IV  
Deep horizontal ablation (raw ENU; one protocol per row)

| Covariance for AR              | Deep $h$ (m)    | Fix rate        |
|--------------------------------|-----------------|-----------------|
| Shadow (paired w/ local study) | $2.77 \pm 0.44$ | $\sim 1.3\%$    |
| Local conditioning (unsafe)    | $6.75 \pm 2.98$ | $\sim 17\%$     |
| Shadow (paired w/ proposed)    | $3.03 \pm 0.68$ | $1.7\%$         |
| Proposed exact-in-window       | $3.40 \pm 0.78$ | $8.8 \pm 1.2\%$ |

Spreads are repeat-run variation on one Deep sequence, not independent scenes. No claim is made that the proposed path reproduces Ceres-fed end-to-end AR decisions.

TABLE V  
Deep system impact: proposed  $Q_{aa}$  vs. shadow (one sequence, 3 repeats)

| Arm      | $h$ (m)         | $u$ (m)         | yaw ( $^\circ$ ) | fix     |
|----------|-----------------|-----------------|------------------|---------|
| Shadow   | $3.03 \pm 0.68$ | $2.27 \pm 0.27$ | $1.09 \pm 0.28$  | $1.7\%$ |
| Proposed | $3.40 \pm 0.78$ | $1.46 \pm 0.21$ | $0.66 \pm 0.07$  | $8.8\%$ |

## VII. Covariance Accuracy and Runtime

Open-sky end-to-end with the proposed path preserves accuracy and fix count inside the frozen band:  $0.028991\text{ m} / 0.4727^\circ$ , 869/1775 fixed versus baseline 876/1775. Thus, when the measurement model is well behaved, joint  $Q_{aa}$  neither inflates fixing nor loses fixes relative to shadow AR.

The local-conditioning approximation is the critical negative control (Table IV): it is real-time but unsafe for AR on Deep. The proposed path removes that catastrophic horizontal regression while keeping a large fraction of the fix-rate increase that joint information enables.

## VIII. End-to-End AR Evaluation

a) System impact of the covariance swap.: Table V reports paired Deep runs that differ in the AR covariance path (proposed vs. shadow), under the same protocol wave. Vertical RMSE and yaw improve in all three repeats; fix rate rises about fivefold; horizontal RMSE is mixed (worse in 2/3 repeats). These repeats are not independent statistical samples; they measure run-to-run variation under load-sensitive replay. Medium is withheld from result tables because its 0% NIFix regime does not isolate AR effects from solver timing and ordering (Sec. IX).

b) Boundary of the end-to-end evidence.: This paper does not claim that proposed AR decisions and trajectories match a Ceres-fed AR oracle. Its primary validation is covariance-level agreement and runtime; UrbanNav results quantify secondary system impact. Fig. 4 shows only that swapping to joint  $Q_{aa}$  can change fix timing and later horizontal error; it is not a causal proof of prior poisoning.

## IX. Limitations

- Exactness scope. Frobenius match at  $10^{-3}$  covers 90.1% of usable epochs. Without Loewner, per-diagonal, and worst-direction tests, the evidence does not establish elementwise equivalence or directional dominance.




Fig. 4. Example Deep horizontal error over time (one proposed-path run with a heavy tail vs. a shadow baseline run). Markers: fixed epochs on the proposed arm.

- No Ceres e2e equivalence claim. Covariance-level agreement does not imply identical AR decisions or trajectories on UrbanNav.
- Medium withheld. At 0% NIFix, arm-to-arm differences can reflect solver timing, WIFix, or ordering rather than the tested AR covariance.
- Soft real-time. The 50 ms figure is a covariance-query budget at  $\sim 1\text{ Hz}$  AR; 4.3% exceed it (max 97.5 ms). It is not a reported bound on total optimize-plus-covariance latency.
- Statistics. Three repeats of one Deep sequence are not three independent environments.
- Decision-layer urban mitigations are outside this letter. In particular, the Deep horizontal heavy tail remains unresolved. Because it is consistent with residual multipath on tracked (line-of-sight) signals rather than covariance-scale error, a GNSS-independent consistency-based float weighting is reserved for a separate follow-on study.

## X. Conclusion

We presented a domain-structured online path to the exact-in-window marginal ambiguity covariance of a tightly coupled GNSS-IMU-vision factor graph, with numerical abstention and shadow fallback. On open-sky graphs it is about  $29\times$  faster than requesting the same blocks from `ceres::Covariance`, produces a validated output on 99.94% of epochs, and achieves Frobenius relative error below  $10^{-3}$  on 90.1% of usable epochs. Directional Loewner diagnostics, a Ceres end-to-end AR oracle, and deterministic float-path isolation would support stronger future claims, but are not assumed by the bounded claims in this paper. Secondary Deep system effects versus shadow AR (more fixes; mixed horizontal) must be interpreted in that light.

## Reproducibility

Opt-in flags, configs, and evaluation scripts are in the accompanying repository (`research/paper/README.md`). Runtime percentiles regenerate from `[vaar-fast]` logs via `scripts/fig_timing.py`. Table generation: `scripts/eval_paper_tables.py`.

## References

- [1] C. Chi, X. Zhang, J. Liu et al., “GICI-LIB: A GNSS/INS/Camera integrated navigation library,” NAVIGATION, 2023, open-source tightly-coupled RTK/IMU/camera; Table V UrbanNav RRR baselines.
- [2] P. J. G. Teunissen, “The least-squares ambiguity decorrelation adjustment: a method for fast GPS integer ambiguity estimation,” *Journal of Geodesy*, vol. 70, no. 1–2, pp. 65–82, 1995.
- [3] X.-W. Chang, X. Yang, and T. Zhou, “MLAMBDA: a modified LAMBDA method for integer least-squares estimation,” *Journal of Geodesy*, vol. 79, no. 9, pp. 552–565, 2005.
- [4] S. Verhagen and P. J. G. Teunissen, “The ratio test for future GNSS ambiguity resolution,” *GPS Solutions*, vol. 17, no. 4, pp. 535–548, 2013.
- [5] H.-J. Euler and B. Schaffrin, “On a measure for the discernibility between different ambiguity solutions in the static-kinematic GPS-mode,” *IAG Symposium*, vol. 107, pp. 285–295, 1991, kinematic Systems in Geodesy, Surveying, and Remote Sensing.
- [6] P. J. G. Teunissen, “Success probability of integer GPS ambiguity rounding and bootstrapping,” *Journal of Geodesy*, vol. 72, no. 10, pp. 606–612, 1998.
- [7] —, “Theory of integer equivariant estimation with application to GNSS,” *Journal of Geodesy*, vol. 77, no. 7–8, pp. 402–410, 2003.
- [8] R. Odolinski and P. J. G. Teunissen, “Best integer equivariant estimation: performance analysis using real data collected under low- to high-ionospheric activity and at low- to high-latitudes,” *Journal of Geodesy*, vol. 94, no. 9, p. 91, 2020.
- [9] T. Qin, P. Li, and S. Shen, “VINS-Mono: A robust and versatile monocular visual-inertial state estimator,” *IEEE Transactions on Robotics*, vol. 34, no. 4, pp. 1004–1020, 2018.
- [10] S. Leutenegger, S. Lynen, M. Bosse, R. Siegwart, and P. Furgale, “Keyframe-based visual-inertial odometry using nonlinear optimization,” in *The International Journal of Robotics Research*, vol. 34, no. 3, 2015, pp. 314–334, oKVIS; FEJ / marginalization consistency.
- [11] S. Cao, X. Lu, and S. Shen, “GVINS: Tightly coupled GNSS–visual–inertial fusion for smooth and consistent state estimation,” *IEEE Transactions on Robotics*, vol. 38, no. 4, pp. 2004–2021, 2022.
- [12] A. M. Erisman and W. F. Tinney, “On computing certain elements of the inverse of a sparse matrix,” *Communications of the ACM*, vol. 18, no. 3, pp. 177–179, 1975.
- [13] L. Lin, J. Lu, L. Ying, R. Car, and W. E, “Fast algorithm for extracting the diagonal of the inverse matrix with application to the electronic structure analysis of metallic systems,” *Communications in Mathematical Sciences*, vol. 7, no. 3, pp. 755–777, 2009, selected/diagonal inversion context; see also sparse selected-inversion literature.
- [14] M. Kaess, H. Johannsson, R. Roberts, V. Ila, J. J. Leonard, and F. Dellaert, “iSAM2: Incremental smoothing and mapping using the Bayes tree,” *The International Journal of Robotics Research*, vol. 31, no. 2, pp. 216–235, 2012.
- [15] F. Dellaert and M. Kaess, “Factor graphs for robot perception,” *Foundations and Trends in Robotics*, vol. 6, no. 1–2, pp. 1–139, 2017.
- [16] S. Agarwal, K. Mierle et al., “Ceres solver,” <http://ceres-solver.org>, 2023.
- [17] P. D. Groves, “Shadow matching: A new GNSS positioning technique for urban canyons,” *Journal of Navigation*, vol. 64, no. 3, pp. 417–430, 2011.
- [18] H.-F. Ng, G. Zhang, and L.-T. Hsu, “A computation effective range-based 3D mapping aided GNSS with NLOS correction method,” *Journal of Navigation*, 2020, 3DMA / NLOS correction literature.
- [19] W. Wen and L.-T. Hsu, “Towards robust GNSS positioning and real-time kinematic using factor graph optimization,” *Proceedings of ION GNSS+*, 2021, nLOS/multipath context in urban canyons.
- [20] L.-T. Hsu et al., “UrbanNav: An open-sourced multisensory dataset for benchmarking positioning algorithms designed for urban areas,” *ION GNSS+*, 2020, iPNL-POLYU UrbanNav-HK Medium/Deep/Harsh.